

## ISA

Components: processor, I/O, mem, and network.

ISA: formal specification of the instruction set that is implemented in the machine hardware.

Simplicity favours regularity; Smaller faster; Make common case fast; Good design demands good compromises;

Important registers:

(ra: usually PC + 4); (sp: must be aligned to 4 bytes);

(gp: holds the base address of global variables)

## Arithmetic

rs1 and rs2 fields need to be kept in the same place: imm field in S-type separated

destination = source1 op source2

In I format, values range:  $-2^{11}$  to  $2^{11} - 1$ .

Load 32 bits:

```
lui t0, 1010 1010 1010 1010 1010b
```

```
ori t0, t0, 1010 1010 1010b
```

logical shift: fill the vacancy with zeros

```
slli t2, s0, 8 # t2 = s0 << 8 bits
```

```
srli t2, s0, 8 # t2 = s0 >> 8 bits
```

```
lw t0, 4(s3) # load word from mem to reg
```

```
sw t0, 8(s3) # store word from reg to mem
```

(loaded or stored using a 5-bit address)

Address is byte-base: increment 4 when accessing reg

Little Endian: rightmost byte is the most significant byte.

lb places the byte from mem into the rightmost 8 bits of the dest reg and signed extension.

```
lb t0, 1(s3) # load byte from memory
```

```
sb t0, 6(s3) # store byte to memory
```

stack grows from high address to low address. in recursive procedures, rd stored in ra, which kept on stack

2's complement: complement all the bits and then add 1

$6 = 00...0110_2 \Rightarrow 11...1001_2 + 1 \Rightarrow 11...1010 = -6$

$n$ -bit signed binary:  $[-2^{n-1}, 2^{n-1} - 1]$

## ALU

32-bit signed numbers: range from  $2^{31} - 1$  to  $-2^{31}$

If the bit string represents address: 0 to  $2^{32} - 1$ .

Sign extension copies the most significant bit into the other bits to preserve the sign of the number.

Ripple Carry Adder: connect all adders in sequence, slow because each bit's carry-out depends on the previous bit's carry-in, leading to a cumulative delay (glitching).

Glitching: invalid and unpredictable output that can be read by the next stage, resulting in incorrect behavior.

Critical path: longest sequence of dependent operations

Overflow: adding two positive numbers yields a negative / adding two negative numbers gives a positive / subtracting a negative from a positive gives a negative / subtracting a positive from a negative gives a positive.

To detect overflow, use carry-in MSB  $\wedge$  carry-out MSB  
 $n$ -bit  $\times$   $m$ -bit multiplication, must have  $n+m$  bits to cover all possible products.

mul: 32-bit  $\times$  32-bit multiplication and places the lower 32 bits in the destination register. mulh, mulhu, and mulhsu perform the same multiplication but return the upper 32 bits of the full 64-bit product.

Logical shifts fill with zeros, while arithmetic right shifts fill with the sign bit.

**Floating**  $\underbrace{6.6254}_{\text{Mantissa (always +)}} \times \underbrace{10}_{\text{Base}}^{-27} \Leftarrow \pm 1.M \times 2^{E'-127}$

Structure: S — E' — M (1 - 8 - 23)

S: Sign bit; E': 8-bit signed exponent; M: mantissa

e.g. 40C0000<sub>16</sub> in decimal

1. 40C0000 = 0 10000001 100000000000000000000000

2. Sign bit (0): Positive (+)

3. Exponent:  $10000001_2 - 127 = 129 - 127 = 2$

4. Mantissa:  $1.1000000000... = 1 + 1 \times 2^{-1} = 1.5$

Result:  $1.5 \times 2^2 = 6$

e.g.  $-0.5_{10}$  in binary

1. Sign bit: 1

2. Mantissa:  $0.5 = 1.0 \times 2^{-1}$

3. Exponent:  $127 - 1 = 126 = 01111110$

Result:  $-0.5_{10} = 10111111000000000000000000000000$

$E = 0, M = 0$ : 0;  $E = 0, M \neq 0$ : denormalized number, which is  $\pm 0.M \times 2^{-126}$ ;  $E = 1...1, M = 0$ :  $\pm\infty$ , depending on the sign;  $E = 1...1, M \neq 0$ : NaN (Not a Number).

## Datapath

combinational: ALU; state: memory

The partition of imm field: align the imm with other instruction types - more efficient implementation of control units.

The instruction is decoded in the path between the Instruction Memory and Register File.

## Pipeline

clock cycle: timed to accommodate the slowest instruction: instr take same amount of time

CPU time = CPI  $\times$  CC  $\times$  IC, CPI = cycles per instruction, CC = clock cycle time, IC = instruction count.

IF: Instruction fetch and PC update

ID: Instruction decode and register file read

EXE: Execution or address calculation

MEM: Data memory access (only in loads and stores)

WB: Write the result data back into the register file  
instruction latency (time from start of instr to completion) is not reduced.

State registers between pipeline stage: flip-flop, data moves in at rising edge

Structural hazards: conflicts in the use of a resource

- separating instruction and data memories

- reads in the 2nd half of the cycle and writes in the 1st half

$$\text{Clock rate} = \frac{1}{\text{Clock cycle time}}$$

Data Hazards: dependence of one instr on an earlier one

- Read After Write (RAW)

- Load-Use data hazard

Solution: 1. Insert NOP / Stall 2. Forwarding

Control Hazards: make decision based on results of 1 instr

1. Unconditional branches: jal, jalr

2. Conditional branches: **beq**, **bne**

3. Exceptions

Solution: 1. Stall

2. delay branching (move branch hardware to ID stage, use delay slot to do other instr first)

3. branch prediction

NOP is determined at the compilation stage, while a flush is determined during runtime

- static branch prediction: assume not taken (works for top of the loop, not for btm of the loop); assume branch taken need one stall cycle

- dynamic branch prediction: use run-time info (branch prediction buffer), if prediction wrong, flush the incorrect instr, restart pipeline

Exception: Traps - division by zero (synchronous); Interrupts - pressing keyboard while compiling (asynchronous)

### Performance

Response Time: start and completion of a task

Throughput: total amount of work done in a given time

$$\text{performance}_X = \frac{1}{\text{execution time}_X}$$

If  $X$  is  $n$  times faster than  $Y$ , then

$$\frac{\text{performance}_X}{\text{performance}_Y} = \frac{\text{execution time}_Y}{\text{execution time}_X} = n$$

CPU exe time = # of CPU clock cycles  $\times$  clock cycle time

$$= \frac{\text{number of CPU clock cycles}}{\text{clock rate}}$$

CPU clock cycles = # of instr  $\times$  clock cycles per instr

$$\text{CPI}_{\text{eff}} = \sum_{i=1}^n \text{CPI}_i \times \text{IC}_i$$

$\text{IC}_i$ : percentage of the no of instructions;  $\text{CPI}_i$ : ave no of clock cycles per instruction for that instruction class.

CPU time = Instruction count  $\times$  CPI  $\times$  clock cycle time

$$\text{GM} = n \cdot \sqrt[n]{\sum_{i=1}^n \text{SPEC ratio}_i}$$

### Memory

1 k  $\approx 2^{10}$  (kilo) 1000; 1 M  $\approx 2^{20}$  (Mega) 1000000;

1 G  $\approx 2^{30}$  1000000000; 1 T  $\approx 2^{40}$  (Tera) 1000000000000

Small memories are fast, large memories are slow.

L1: Cache using Static RAM (SRAM)

L2: Another cache, using faster Dynamic RAM (DRAM)

L3: Main memory, typically DRAM

L4: Secondary memory such as flash storage or solid-state drives (SSD)

first three: volatile - retain data when power on

Temporal Locality (time), if referenced, will be referenced again soon: keeps the most recently accessed data closer to processor; Spatial (space - array), location is referenced, locations with nearby addresses will be referenced soon

RAM: SRAM, DRAM, SDRAM, and DDR SDRAM

SRAM:  $\geq 6$  transistors, fast, large, expensive

DRAM: 1 transistor, capacitor, slow, small, cheap

Interleaving: hide memory access latency

Secondary memory: e.g. magnetic disk

Reg - Cache (SRAM) - Main Mem (DRAM) - Disk - Tape

### Cache

Direct mapping: tag (identify which block from MM) + blk no (index the cache) + byte address

Direct mapped cache with  $2^n$  blocks,  $n$  bits are used for the index. For a block size of  $2^m$  words ( $2^{m+2}$  bytes),  $m$  bits are used to address the word within the block.

2 bits are used to address the byte within the word.

Tag size = 32 - ( $n + m + 2$ )

The total number of bits in a direct-mapped cache is:

$2^n \times (\text{block size} + \text{tag field size} + \text{valid field size})$

1 word = 4 bytes = 32 bits (diff bit hv diff calculation)

Associative Mapping: arbitrary block

Combine: use tag + set + byte

cache write hit: write-through: write both mem and cache (use buffer); write-back: only write cache and write to mem when cache is replaced (dirty bit)

Asso replace: least recently used / random replace

Average Memory Access Time =  $h \times C + (1 - h) \times M$

$h$  - hit rate,  $C$  - cache access time,  $M$  - miss penalty

Miss Penalty: total access time experienced by the processor when a miss occurs ( $1 + X + 1$ )

high-performance,  $C_1$ : cache access time,  $C_2$ : miss penalty from L2 to L1,  $M$  is from main memory to L2 to L1:  $h_1 \times C_1 + (1 - h_1) \times [h_2 \times C_2 + (1 - h_2) \times M]$

### Virtual Memory

Address space: pages (fixed size) or segments (variable sizes). Frequently used blocks are copied into the cache.

page table: stores mapping between virtual and physical pages. recent translations may be cached in the Translation Lookaside Buffer for faster access.

### Instruction Level Parallelism

Multiple-Issue: instructions per cycle (IPC)

static multiple-issue processor (VLIW): compiler decides which instructions can run in parallel; dynamic multiple-issue processor (superscalar): CPU decides at runtime  
Data hazards: True Dependency (RAW); Anti-dependency (WAR); Output Dependency (WAW).

Register renaming: rename original reg to a new reg

VLIW: (1) instruction scheduling: must separate load-use instrs from their loads by one cycle, avoid stalls, and (2) loop unrolling (copy the loop body to reduce loop)

Superscalar: Dynamic multiple-issue processors decide at runtime which instrs to run in parallel