

CSCI3160 Design and Analysis of Algorithms

Ryan Chan

November 2, 2025

Abstract

This is a note for **CSCI3160 Design and Analysis of Algorithms**.

Contents are adapted from the lecture notes of CSCI3160, prepared by [Xiao Liang](#) and [Yufei Tao](#), as well as some online resources.

This note is intended solely as a study aid. While I have done my best to ensure the accuracy of the content, I do not take responsibility for any errors or inaccuracies that may be present. Please use the material thoughtfully and at your own discretion.

If you believe any part of this content infringes on copyright, feel free to contact me, and I will address it promptly.

Mistakes might be found. So please feel free to point out any mistakes.

Contents

1	Introduction	2
1.1	The RAM Model	2
1.2	Efficiency of the Worst Input	4
1.3	Basic Techniques	5
2	Divide and Conquer	8
2.1	Sorting	8
2.2	Counting Inversions	8
2.3	Dominance Counting	9
2.4	Matrix Multiplication (Strassen's Algorithm)	9
2.5	Fast Fourier Transform	10
3	Greedy Algorithms	14
3.1	Activity Selection	14
4	Dynamic Programming	16
5	Graph	17
A	Master Theorem	18

Chapter 1

Introduction

Computer science is a subject where we first define a computation model, which is a simple yet accurate abstraction of a computing machine, and then we gradually build up a theory based on this model.

Thus, the first thing to do is to come up with a computation model. For the study of algorithms, we need to consider the following criteria:

1. **Mechanically Implementable:** it can run real algorithms.
2. **Sufficiently General:** it can capture all the natural steps of problem-solving.
3. **Sensitive Enough:** it can distinguish differences in resource usage (time and space).

We choose the **Random Access Machine (RAM)** as our model.

1.1 The RAM Model

The RAM model has a **memory** and a **CPU**.

Memory is defined as an infinite sequence of cells, each containing the same number of w bits, with addresses 1, 2, and so on. The CPU contains a fixed number of registers, each of which has w bits. A word is a sequence of w bits, where w is called the word length. In other words, each memory cell and CPU register stores a word.

1.1.1 Atomic Operations

We say that there are a few atomic operations that the CPU can perform.

1. Register (Re-)Initialization

We can set a register to a fixed value or to the content of another register.

2. Arithmetic

This operation takes two integers stored in two registers and performs basic arithmetic calculations.

3. Comparison / Branching

This operation takes two integers stored in two registers, compares them, and determines the result of the comparison.

4. Memory Access

This operation takes a memory address that is currently stored in a register, then performs either reading (to register) or writing (to memory).

5. Randomness

`RANDOM(x, y)` returns an integer chosen **uniformly** at random in $[x, y]$, where $x \leq y$. The resulting random integer is then placed in a register.

An execution is defined as a sequence of atomic operations, where the cost (running time) of an execution is the length of such a sequence, i.e., the number of atomic operations it performs.

1.1.2 Algorithms

We first take a look at some terminologies.

An input refers to the initial state of the registers and the memory before an execution starts.

An algorithm is a description that, given an input, can be utilized to **unambiguously** produce a sequence of atomic operations — namely, the execution of the algorithm. In other words, it should always be clear what the next atomic operation should be, given the outcomes of all the previous atomic operations.

The cost of an algorithm on an input is the length of its execution on that input (i.e., the number of atomic operations required).

The space of an algorithm on an input is the largest memory address accessed by the algorithm's execution on that input.

We define an algorithm as **deterministic** if it never invokes the atomic operation `RANDOM`; otherwise, the algorithm is **randomized**.

1.1.3 Expected Running Time

A deterministic algorithm has a fixed cost for the same input, whereas for a randomized algorithm, the cost is a random variable, since for each input the cost might change every time the algorithm is executed.

Algorithm 1.1: Find a One

```
1 while r ≠ 1 do
2   r = RANDOM(0, 1)
3 return r = 1
```

Here, we cannot know how many times line 2 will be executed, as each execution can produce a new sequence of atomic operations.

Thus, we use the **expected cost** to evaluate the cost of a randomized algorithm.

Definition 1.1.1. Let X be a random variable that represents the cost of an algorithm on a given input. The expected cost of the algorithm on that input is the expectation of X .

We use the expected running time, rather than other metrics, for the following reasons:

1. Ease of computation
2. Linearity of expectation: it allows us to break the analysis into smaller parts
3. Concentration bounds

Theorem 1.1.1 (Markov's Inequality). Suppose that a random variable $X \geq 0$ only takes non-negative values. Then, for every $t > 0$,

$$\mathbb{P}(X \geq t) \leq \frac{\mathbb{E}[X]}{t}.$$

The theorem here shows that if the average value of X is small, then it is less likely that X is large. This provides an upper bound on the probability that a non-negative random variable is much larger than its expectation.

Let T be the random variable representing the running time of a randomized algorithm, and suppose $T \geq 0$ and $\mathbb{E}[T] = \mu$. Then, for any $c > 1$,

$$\mathbb{P}(T \geq c\mu) \leq \frac{1}{c}.$$

Example (Las Vegas Algorithm). A Las Vegas algorithm always gives the correct result but has a random running time.

Suppose $\mathbb{E}[T] = \mu$, and we run this algorithm with a timeout of $c\mu$. Then we have

$$\mathbb{P}[\text{timeout}] \leq \frac{1}{c}.$$

This shows that if we run the algorithm for c times its expected running time, then the probability that it fails will be less than or equal to $\frac{1}{c}$.

For example, if we run for 3μ steps, then it will fail with probability $\leq \frac{1}{3}$. Therefore, if we repeat it 3 times, we will succeed with probability $1 - (\frac{1}{3})^3 = 0.963$.

1.2 Efficiency of the Worst Input

By looking at the cost of an algorithm on the worst input, we are able to measure the speed of an algorithm.

Definition 1.2.1. Define \mathcal{I}_n , where n is an integer, to be the set of all inputs to a problem that have the same problem size n . Given an input $I \in \mathcal{I}_n$, the cost $X_{\mathcal{A}}(I)$ of an algorithm \mathcal{A} is the length of its execution on I .

The worst-case cost of \mathcal{A} under the problem size n is the maximum $X_{\mathcal{A}}(I)$ over all $I \in \mathcal{I}_n$.

The worst expected cost of \mathcal{A} under the problem size n is the maximum $\mathbb{E}[X_{\mathcal{A}}(I)]$ over all $I \in \mathcal{I}_n$.

Example (Dictionary Search). In memory, a set S of n integers has been arranged in ascending order in the memory cells from address 1 to n . The value of n has been placed in Register 1 of the CPU. Another integer v has been placed in Register 2, where n is the problem size, and \mathcal{I}_n is the set of all possible (S, v) .

We want to determine whether v exists in S .

The worst-case cost of the [binary search algorithm](#) is $\mathcal{O}(\log n)$. This means that on any input in \mathcal{I}_n , the maximum number $f(n)$ of atomic operations performed by the algorithm grows no faster than $\log_2 n$.

Example (Randomized Algorithm). Consider the following randomized algorithm:

Algorithm 1.2: Find a Zero

Data: A is an array of size n that contains at least one 0

```

1 while  $A[r] \neq 0$  do
2    $r = \text{RANDOM}(1, n)$ 
3 return  $r$ 

```

The expected cost depends on the input array. For example, if all numbers in A are 0, then the algorithm finishes in $\mathcal{O}(1)$. If it has only one 0, the algorithm will finish in $\mathcal{O}(n)$ because each $A[r]$ has a $\frac{1}{n}$ probability of being 0, and we need to repeat n times in expectation to find the 0.

Thus, we have worst-case cost = ∞ and worst expected cost = $\mathcal{O}(n)$.

Example (Find a Zero). Let A be an array of n integers, among which half are 0. Design an algorithm to report an arbitrary position of A that contains a 0.

Algorithm 1.3: Find a Zero

```
1 while  $A[r] \neq 0$  do
2    $r = \text{RANDOM}(1, n)$ 
3 return  $r$ 
```

The algorithm finishes in $\mathcal{O}(1)$ expected time on every input A .

Proof. Let X be the number of iterations until the algorithm picks a zero. Each iteration chooses $r \in \{1, \dots, n\}$.

Then we have

$$p = \mathbb{P}[A[r] = 0] = \frac{\# \text{ of zeros}}{n} = \frac{\frac{n}{2}}{n} = \frac{1}{2}.$$

Then X is a geometric random variable with success probability $p = \frac{1}{2}$.

For a geometric random variable, we have

$$\mathbb{E}[X] = \frac{1}{p} = \frac{1}{\frac{1}{2}} = 2.$$

Since each iteration takes $\mathcal{O}(1)$ time to pick a random index, we need two such iterations in expectation. Thus, the expected running time is $2 \times \mathcal{O}(1) = \mathcal{O}(1)$. \blacksquare

Remark. In contrast, any deterministic algorithm must probe at least $\frac{n}{2}$ integers of A in the worst case. In other words, any deterministic algorithm must have a worst-case time of $\Theta(n)$ — provably slower than the above randomized algorithm (in expectation).

1.3 Basic Techniques

In algorithm, there are three basic techniques we can utilize.

1.3.1 Recursion

When dealing with a sub-problem (the same problem but with a smaller input), consider it solved, and use the sub-problem's output to continue the algorithm design.

Example (The Hanoi Tower Problem). There are 3 rods A, B, and C. On rod A, n disks of different sizes are stacked, such that no disk of a larger size is above a disk of a smaller size. The other two rods are empty.

It is allowed to move the top-most disk of a rod to another, while no disk of a larger size can be put above a disk of a smaller size. The goal is to design an algorithm to move all the disks to rod B.

To move the largest disk, i.e., the n -th disk, to rod B, we first need to move all disks above it to rod C. For all disks, we use this method; thus, if we ignore the last disk, the remaining problem is to move $n - 1$ disks. This is the same problem with a smaller size.

Suppose that the algorithm performs $f(n)$ operations to solve a problem of size n , where $f(1) = 1$. By recursion, we have

$$f(n) \geq 1 + 2 \times f(n - 1).$$

Solving this gives

$$f(n) \geq 2^n - 1.$$

Thus, the best time complexity for solving this problem is $\Omega(2^n)$, where n is the number of disks.

1.3.2 Repeating till Success

Given a set S of n integers in an array and an integer $k \in [1, n]$, we want to find the k -th smallest integer of S .

Definition 1.3.1 (Rank). The rank of an integer $v \in S$ is the number of elements in S smaller than or equal to v .

For example, suppose that $S = (53, 92, 85, 23, 35, 12, 68, 74)$, then the rank of 53 is 4.

We can obtain the rank of v in $\mathcal{O}(|S|)$ time.

Example (Sub-problem of The k -Selection Problem). Assume n is a multiple of 3. We want to obtain a subproblem of size at most $\frac{2n}{3}$ with exactly the same result as the original problem.

Our goal is to produce a set S' and an integer k' such that $|S'| \leq \frac{2n}{3}$, $k' \in [1, |S'|]$, and the element with rank k' in S' is the element with rank k in S , where it is possible that $k \neq k'$.

Consider the following algorithm A_{sub} :

We first take an element $v \in S$ uniformly at random. Then we divide S into S_1 and S_2 , where S_1 is the set of elements in S that are less than or equal to v , and S_2 is the set of elements in S greater than v .

If $|S_1| \geq k$, then return $S' = S_1$ and $k' = k$. Otherwise, return $S' = S_2$ and $k' = k - |S_1|$.

This algorithm succeeds if $|S'| \leq \frac{2n}{3}$, and fails otherwise.

We repeat this algorithm until it succeeds.

Lemma 1.3.1. The algorithm succeeds with probability at least $\frac{1}{3}$.

Proof. Let S be a set of size n and let $v \in S$ be chosen uniformly at random. Let the rank of v in S be $\text{rank}(v)$.

Define $S_1 = \{x \in S : x \leq v\}$, $S_2 = \{x \in S : x > v\}$. If $|S_1| \geq k$, return $S' = S_1$ and $k' = k$; otherwise $S' = S_2$ and $k' = k - |S_1|$. The sub-problem succeeds if $|S'| \leq \frac{2n}{3}$.

If $|S_1| \geq k$, then $|S'| = |S_1| = \text{rank}(v)$, and success requires

$$\text{rank}(v) \leq \frac{2n}{3}.$$

If $|S_1| < k$, then $|S'| = |S_2| = n - \text{rank}(v)$, and success requires

$$n - \text{rank}(v) \leq \frac{2n}{3} \implies \text{rank}(v) \geq \frac{n}{3}.$$

Then we have

$$\text{rank}(v) \in \left[\frac{n}{3}, \frac{2n}{3} \right].$$

Since v is chosen uniformly at random, every rank is equally likely, and the number of ranks in the success range is at least $\frac{2n}{3} - \frac{n}{3} = \frac{n}{3}$. Thus we have

$$\mathbb{P}[\text{success}] \geq \frac{\frac{n}{3}}{n} = \frac{1}{3}.$$

■

Remark. Choosing $\frac{2n}{3}$ is not mandatory, but rather a convenient threshold since it gives a smaller sub-problem. For example, if we say it succeeds when $|S'| \leq \frac{9n}{10}$, we have the range $[\frac{n}{10}, \frac{9n}{10}]$, giving the probability of success $\frac{8}{10}$, which is high. However, this is simply because we have a larger range to choose from, making the sub-problem larger.

In general, if an algorithm succeeds with a probability at least $c > 0$, then the number of repeats needed

for the algorithm to succeed for the first time is at most $\frac{1}{c}$ in expectation.

1.3.3 Geometric Series

A geometric sequence is an infinite sequence of the form

$$n, cn, c^2n, c^3n, \dots$$

where n is a positive number and c is a constant satisfying $0 < c < 1$. It holds in general that

$$\sum_{i=0}^{\infty} c^i n = \lim_{i \rightarrow \infty} n \cdot \frac{1 - c^i}{1 - c} = \frac{n}{1 - c} = \mathcal{O}(n).$$

The summation $\sum_{i=0}^{\infty} c^i n$ is called a geometric series.

Consider again the k -th selection problem. Using the previous repeating technique, we can convert the problem into a subproblem with size at most $\lceil \frac{2n}{3} \rceil$ in $\mathcal{O}(n)$ expected time. We can use the geometric series to find the expected running time:

$$\begin{aligned} a \cdot n + a \cdot \frac{2}{3} \cdot n + a \cdot \left(\frac{2}{3}\right)^2 \cdot n + \dots &= a \cdot n + a \cdot \sum_{i=1}^{\infty} \left(\frac{2}{3}\right)^i \cdot n \\ &= a \cdot n + a \cdot \mathcal{O}(n) \\ &= \mathcal{O}(n). \end{aligned}$$

Next, we analyze the running time of A_{sub} . It takes $\mathcal{O}(1)$ to select the element v , and dividing S into S_1 and S_2 takes $\mathcal{O}(n)$. Comparing $|S_1|, |S_2|$, and k also takes $\mathcal{O}(n)$. Thus, by the linearity of expectation, every conversion takes $\mathcal{O}(n)$ running time in expectation.

By the previous lemma, we need to repeat the algorithm 3 times in expectation until it succeeds, and each execution takes $\mathcal{O}(n)$ time. Therefore, the k -th selection algorithm takes $\mathcal{O}(n)$ time in expectation.

Chapter 2

Divide and Conquer

We take a look at divide and conquer, which guarantees strong performance. When dividing, we utilize recursion to reduce the problem into sub-problems, and the conquer part tackles the original problem.

2.1 Sorting

We consider [merge sort](#) in this example.

Problem Statement

Given an array A of n distinct integers, produce another array where the same integers are arranged in ascending order.

Divide

Let A_1 be the array containing the first $\lceil \frac{n}{2} \rceil$ elements of A , and A_2 be the array containing the remaining elements of A . We sort A_1 and A_2 recursively.

Conquer

Merge the two sorted arrays in ascending order, which can be done in $\mathcal{O}(n)$ time.

Analysis

Let $f(n)$ denote the worst-case cost of the algorithm on an array of size n . Then,

$$f(n) \leq 2f\left(\left\lceil \frac{n}{2} \right\rceil\right) + \mathcal{O}(n),$$

which gives¹

$$f(n) = \mathcal{O}(n \log n).$$

2.2 Counting Inversions

Problem Statement

Given an array A of n distinct integers, count the number of inversions, where an inversion is a pair (i, j) such that $1 \leq i < j \leq n$ and $A[i] > A[j]$.

Divide

Let A_1 be the array containing the first $\lceil \frac{n}{2} \rceil$ elements of A , and A_2 be the array containing the remaining elements of A . We solve the problem recursively on A_1 and A_2 .

Conquer

It remains to count the number of *crossing inversions* (i, j) where $i \in A_1$ and $j \in A_2$. Using merge sort, we can sort A_1 in $\mathcal{O}(n \log n)$ time. For each element $e \in A_2$, we count how many crossing inversions

¹See Master Theorem [A](#).

e produces using binary search. In total, there are $\frac{n}{2}$ binary searches performed, each taking $\mathcal{O}(\log n)$ time, giving $\mathcal{O}(n \log n)$ time for this step.

Analysis

A simple comparison of each pair takes

$$(n-1) + (n-2) + \cdots + 1 = \mathcal{O}(n^2).$$

Let $f(n)$ denote the worst-case cost of the algorithm on an array of size n . Then,

$$f(n) \leq 2f\left(\left\lceil \frac{n}{2} \right\rceil\right) + \mathcal{O}(n \log n),$$

which gives

$$f(n) = \mathcal{O}(n \log^2 n).$$

2.3 Dominance Counting

Problem Statement

Denote \mathbb{Z} as the set of integers. Given a point p in \mathbb{Z}^2 , denoted by $p[1], p[2]$ its x - and y -coordinate, given two distinct points p and q , we say that q dominates p if $p[1] \leq q[1]$ and $p[2] \leq q[2]$.

Let P be a set of n points in \mathbb{Z}^2 with distinct x -coordinates. Find for each point $p \in P$ the number of points in P that are dominated by p .

Assume that, without loss of generality, the points are given in ascending order in the x -coordinates, i.e.

$$p_1[1] < p_2[1] < \cdots < p_n[1].$$

Divide

Let l be the vertical line such that P has $\lceil \frac{n}{2} \rceil$ points on each side of the line, where P_1 is the set of points of P on the left of l , and P_2 is the set of points on the right of l . We solve the problem recursively on P_1 and P_2 .

Conquer

It remains to count for each point $p_2 \in P_2$ how many points in P_1 it dominates. We sort P_1 by its y -coordinate. Then for each point $p_2 \in P_2$, we obtain the number of points in P_1 dominated by p_2 using binary search.

Analysis

Let $f(n)$ denote the worst-case cost of the algorithm on n points. In total, there are $\frac{n}{2}$ binary searches performed, each taking $\mathcal{O}(\log n)$ time, giving $\mathcal{O}(n \log n)$ time for this step. Then,

$$f(n) \leq 2f\left(\left\lceil \frac{n}{2} \right\rceil\right) + \mathcal{O}(n \log n),$$

which gives

$$f(n) = \mathcal{O}(n \log^2 n).$$

2.4 Matrix Multiplication (Strassen's Algorithm)

Given two $n \times n$ matrices A and B , compute their product AB . Assume for simplicity that n is a power of 2. We can divide each of A and B into 4 sub-matrices of order $\frac{n}{2}$.

Suppose we want to compute

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}.$$

We need to perform 8 order- $\frac{n}{2}$ matrix multiplications in the most trivial case, i.e.

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \end{bmatrix}.$$

In the trivial case, we need $\mathcal{O}(n^3)$ time.

In a non-trivial case, we have

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} p_5 + p_4 - p_2 + p_6 & p_1 + p_2 \\ p_3 + p_4 & p_1 + p_5 - p_3 - p_7 \end{bmatrix}$$

where

$$\begin{aligned} p_1 &= A_{11}(B_{12} - B_{22}) \\ p_2 &= (A_{11} + A_{12})B_{22} \\ p_3 &= (A_{21} + A_{22})B_{11} \\ p_4 &= A_{22}(B_{21} - B_{11}) \\ p_5 &= (A_{11} + A_{22})(B_{11} + B_{22}) \\ p_6 &= (A_{12} - A_{22})(B_{21} + B_{22}) \\ p_7 &= (A_{11} - A_{21})(B_{11} + B_{12}) \end{aligned}$$

If $f(n)$ is the worst-case time of computing the product of two order- n matrices, then each of $p_i, 1 \leq i \leq 7$ can be computed in $f(\frac{n}{2}) + \mathcal{O}(n^2)$ time, where $f(\frac{n}{2})$ is exactly one multiplication between two order- $\frac{n}{2}$ matrices, and $\mathcal{O}(n^2)$ is the cost of matrix additions and subtractions.

Therefore,

$$f(n) \leq 7f\left(\frac{n}{2}\right) + \mathcal{O}(n^2)$$

can be solved to

$$f(n) = \mathcal{O}(n^{\log_2 7}) = \mathcal{O}(n^{2.81}).$$

Remark. For n being a power of 2, we can recursively split the matrix into sub-matrices, and it still takes $\mathcal{O}(n^{2.81})$. This also works for any n since we can pad with zeros if n is not a power of two. Given that the logic here does not depend on n , the running time remains $\mathcal{O}(n^{2.81})$.

2.5 Fast Fourier Transform

For degree- d polynomials,

$$A(x) = \sum_{i=0}^d a_i x^i, \quad B(x) = \sum_{i=0}^d b_i x^i,$$

their product $C(x) = A(x)B(x) = \sum_{k=0}^{2d} c_k x^k$ with

$$c_k = \sum_{i=0}^k a_i b_{k-i}$$

where $a_i, b_i = 0$ if $i > d$.

To do such multiplication it takes $\Theta(d^2)$. To speed up, we again use divide and conquer.

2.5.1 Point-wise Multiplication

First, we take a look at the representation of a polynomial. Consider two representations of $A(x)$:

1. Coefficients a_0, a_1, \dots, a_d
2. Values at $d + 1$ distinct points (point-value pairs):

$$(x_0, A(x_0)), (x_1, A(x_1)), \dots, (x_d, A(x_d))$$

Remark. A degree- d polynomial is uniquely determined by its values at any $d + 1$ distinct points.

If we use point-value representation, given $C(x) = A(x)B(x)$, for any point z we have $C(z) = A(z)B(z)$. Thus, it takes linear time to compute all the values of $C(x)$, i.e. $(2d + 1) = \mathcal{O}(d)$ multiplications.

2.5.2 Root of Unity

Before obtaining $C(x)$, we need to evaluate $A(x)$ and $B(x)$. Evaluating a degree- n polynomial at one point costs at least $\mathcal{O}(n)$ time; since there are n points in total, this takes $\Theta(n^2)$. However, consider the following.

A polynomial can be decomposed as (assume n is even):

$$\begin{aligned} A(x) &= a_0 + a_1x + a_2x^2 + \cdots + a_nx^n \\ &= (a_0 + a_2x^2 + \cdots + a_nx^n) + x(a_1 + a_3x^2 + \cdots + a_{n-1}x^{n-2}) \\ &= A_{\text{even}}(x^2) + x \cdot A_{\text{odd}}(x^2) \end{aligned}$$

where

$$A_{\text{even}}(x) = a_0 + a_2x + \cdots + a_nx^{\frac{n}{2}}, \quad A_{\text{odd}}(x) = a_1 + a_3x + \cdots + a_{n-1}x^{\frac{n}{2}-1}.$$

Then, if we consider paired points $\pm x_i$, we have

$$A(x_i) = A_{\text{even}}(x_i^2) + x_i A_{\text{odd}}(x_i^2), \quad A(-x_i) = A_{\text{even}}(x_i^2) - x_i A_{\text{odd}}(x_i^2).$$

Then the evaluation on points is reduced to the evaluation of $\{x_0^2, \dots, x_{\frac{n}{2}-1}^2\}$ plus a linear-time combination.

However, this reduction cannot utilize recursion; therefore, we choose another set of points using the concept of complex numbers.

Theorem 2.5.1 (Fundamental Theorem of Algebra). Every polynomial of degree $n \geq 1$ with complex coefficients has exactly n roots in \mathbb{C} , counted with multiplicity. Equivalently, for

$$p(z) = a_n z^n + a_{n-1} z^{n-1} + \cdots + a_1 z + a_0, \quad (a_n \neq 0),$$

there exist $\zeta_1, \dots, \zeta_n \in \mathbb{C}$ such that

$$p(z) = a_n \prod_{k=1}^n (z - \zeta_k),$$

where each root appears according to its multiplicity.

Explanation. Consider our problem: we have a polynomial $C(x)$ that can be determined by $n = 2d + 1$ points. According to the theorem, every polynomial of degree $2d$ has exactly $2d$ roots in \mathbb{C} . Hence, we can use these complex roots to find a convenient set of points that make evaluation easier. ⊗

Definition 2.5.1. An n -th root of unity, where n is a positive integer, is a complex number z satisfying

$$z^n = 1.$$

There are n distinct n -th roots of unity, given by

$$z_k = \cos\left(\frac{2k\pi}{n}\right) + i \sin\left(\frac{2k\pi}{n}\right), \quad \forall k \in \{0, 1, \dots, n-1\}.$$

Let

$$\omega_n = \cos\left(\frac{2\pi}{n}\right) + i \sin\left(\frac{2\pi}{n}\right),$$

then we have

$$\begin{aligned}\omega_n^0 &= \cos(0) + i \sin(0) = 1, \\ \omega_n^1 &= \cos\left(\frac{2\pi}{n}\right) + i \sin\left(\frac{2\pi}{n}\right), \\ \omega_n^2 &= \cos\left(\frac{4\pi}{n}\right) + i \sin\left(\frac{4\pi}{n}\right), \\ &\vdots \\ \omega_n^{n-1} &= \cos\left(\frac{2(n-1)\pi}{n}\right) + i \sin\left(\frac{2(n-1)\pi}{n}\right).\end{aligned}$$

Thus, we can write the n -th roots of unity as

$$\omega_n^0, \omega_n^1, \omega_n^2, \dots, \omega_n^{n-2}, \omega_n^{n-1}.$$

By Euler's formula,

$$e^{i\theta} = \cos(\theta) + i \sin(\theta),$$

we have

$$\omega_n^k = \cos\left(\frac{2k\pi}{n}\right) + i \sin\left(\frac{2k\pi}{n}\right) = e^{i\frac{2k\pi}{n}}.$$

Moreover, we note some important facts about the roots of unity:

Proposition 2.5.1. For all $n \in \mathbb{N}$, there are exactly n distinct n -th roots of unity, i.e.,

$$U_n = \{\omega_n^0, \omega_n^1, \omega_n^2, \dots, \omega_n^{n-1}\}.$$

Proposition 2.5.2. For all even $n \in \mathbb{N}$, it holds that

$$\omega_n^{k+\frac{n}{2}} = -\omega_n^k, \quad \forall k \in \{0, 1, \dots, \frac{n}{2} - 1\}.$$

Proposition 2.5.3. 3. For all even $n \in \mathbb{N}$, squaring each element of U_n gives the set $U_{\frac{n}{2}}$, i.e., the set of $\frac{n}{2}$ -th roots of unity. More explicitly,

$$\omega_n^{2k} = \omega_{\frac{n}{2}}^{k \bmod (\frac{n}{2})}.$$

Proof. Let $\omega_n = e^{\frac{2\pi i}{n}}$ be a primitive n -th root of unity. Consider the square of ω_n^k :

$$(\omega_n^k)^2 = (e^{\frac{2k\pi i}{n}})^2 = e^{\frac{2(2k)\pi i}{n}} = e^{\frac{2k\pi i}{\frac{n}{2}}} = \omega_{\frac{n}{2}}^k.$$

Thus, for each $k = 0, 1, \dots, n-1$, squaring ω_n^k gives $\omega_{\frac{n}{2}}^k$. ■

Remark.

- Since k ranges over $0, \dots, n-1$, some values of $\omega_{\frac{n}{2}}^k$ repeat (because k modulo $\frac{n}{2}$).
- Therefore, the set of all $(\omega_n^k)^2$ is exactly

$$\{\omega_{\frac{n}{2}}^0, \omega_{\frac{n}{2}}^1, \dots, \omega_{\frac{n}{2}}^{n/2-1}\} = U_{\frac{n}{2}},$$

the set of $\frac{n}{2}$ -th roots of unity.

2.5.3 Fast Fourier Transform

Now we can look at the details of the Fast Fourier Transform (FFT).

Given a polynomial $A(x)$ of degree at most $n-1$, where $n = 2^m$ for some $m \geq 0$.

If $n = 1$, then $\deg A(x) \leq 0$, and $A(x) = a_0$.

If $n \geq 2$,

1. Split $A(x)$ into its even and odd parts.
2. Make two recursive FFT calls of size $\frac{n}{2}$ to evaluate both $A_{\text{even}}(y)$ and $A_{\text{odd}}(y)$ on all the points in

$$U_{\frac{n}{2}} = \{\omega_n^0, \omega_n^1, \dots, \omega_n^{\frac{n}{2}-1}\}.$$

3. For all $k \in \{0, 1, \dots, n-1\}$, compute

$$A(\omega_n^k) = A_{\text{even}}(\omega_n^{2k}) + \omega_n^k A_{\text{odd}}(\omega_n^{2k}),$$

using the values of $A_{\text{even}}(y)$ and $A_{\text{odd}}(y)$ at $\omega_n^{2k} = \omega_{\frac{n}{2}}^{k \bmod (\frac{n}{2})}$.

4. Output the values $(A(\omega_n^k))_{k=0}^{n-1}$.

Then we have the time complexity

$$f(n) = \begin{cases} \mathcal{O}(1), & \text{if } n = 1; \\ 2f\left(\frac{n}{2}\right) + \mathcal{O}(n), & \text{if } n \geq 2. \end{cases}$$

Explanation. Suppose we want to evaluate a polynomial at $n = 8$ points.

1. Split into even and odd parts:

$$A(x) = A_{\text{even}}(x^2) + x A_{\text{odd}}(x^2)$$

This reduces the 8-point evaluation to two 4-point evaluations at x^2 .

2. Recursively split each 4-point evaluation into even/odd again, reducing to 2-point evaluations at x^4 .
3. Split again until reaching 1-point evaluations (just the coefficients themselves).
4. Combine results on the way back using

$$A(\omega_n^k) = A_{\text{even}}(\omega_{n/2}^k) + \omega_n^k A_{\text{odd}}(\omega_{n/2}^k),$$

which takes $\mathcal{O}(n)$ time per level.

Key idea: At each level, the recursive results are reused for multiple points; only the multiplication by ω_n^k differs. This divide-and-conquer gives total complexity $\mathcal{O}(n \log n)$. ⊗

The last step is to turn the point-value pair representation back to the coefficient representation, which is called interpolation. We can accomplish the interpolation in $\mathcal{O}(n \log n)$ time by using the inverse FFT.

2.5.4 Analysis

We can summarize all the operations here.

First, we select the smallest $n \geq 2d + 1$ where $n = 2^m$, $m \geq 0$.

Then we do the evaluation using FFT for both $A(\omega_n^k)$ and $B(\omega_n^k)$, which takes $\mathcal{O}(n \log n)$.

We then do the point-wise multiplication, i.e. $C(\omega_n^k) = A(\omega_n^k)B(\omega_n^k)$, which takes $\mathcal{O}(n)$.

Last, we do the interpolation using inverse FFT, taking $\mathcal{O}(n \log n)$ time.

Thus, we can finish a degree- n polynomial multiplication in time $\mathcal{O}(n \log n)$.

Chapter 3

Greedy Algorithms

Greedy algorithms enforce a simple strategy, i.e. make the locally optimal decision at each step.

3.1 Activity Selection

Problem Statement

Given a set S of n intervals of the form $[s, f]$ where s, f are integers, we want to output a subset $T \subseteq S$ of disjoint intervals with the largest size $|T|$.

Remark. One can think of $[s, f]$ as the duration of an activity, and consider the problem as picking the largest number of activities that do not have time conflicts.

Greedy Algorithm

The idea is to pick the earliest-finishing interval, keep it, and discard all intervals that overlap it. Repeat until none remain.

Algorithm:

Repeat until S becomes empty:

- Add to T the interval $\mathcal{I} \in S$ with the smallest finish time.
- Remove from S all intervals intersecting \mathcal{I} (including \mathcal{I}).

The greedy algorithm picks the earliest finishing activity, leaving the most room for future choices.

Proof. Let $G = \{g_1, g_2, \dots, g_k\}$ be the greedy solution, and $O = \{o_1, o_2, \dots, o_m\}$ be an arbitrary optimal solution, where each g_i, o_i is an interval.

Since there is no overlap, it holds that for each i , $f_{o_{i-1}} \leq s_{o_i}$. It must be $k \leq m$; otherwise, O would not be an optimal solution.

Assume g_1 is the earliest finishing interval, as well as o_1 , which might be different. Then we have $f_{g_1} \leq f_{o_1}$.

We can replace o_1 with g_1 in O . Since g_1 finishes no later than o_1 , intervals in O remain disjoint.

Repeating this process, assume $k < m$. The replacement stops at

$$O = \{g_1, g_2, \dots, g_k, o_{k+1}, o_{k+2}, \dots, o_m\}.$$

Note that the finish time of g_k is earlier than the start time of all intervals $o_{k+1}, o_{k+2}, \dots, o_m$, where disjointness still holds. Then the greedy algorithm would not stop at g_k , because there are still intervals remaining in the original set S . This leads to a contradiction. Thus, it must be $k = m$, i.e. the greedy algorithm gives an optimal solution. ■

Analysis

This algorithm can be implemented in $\mathcal{O}(n \log n)$ time.

Chapter 4

Dynamic Programming

Chapter 5

Graph

Appendix A

Master Theorem

We can use the Master Theorem to solve recurrence problems.

Consider

$$T(n) = aT\left(\frac{n}{b}\right) + f(n), \quad a \geq 1, \quad b > 1,$$

where

$$f(n) = \mathcal{O}(n^k \log^p n).$$

Case 1: if $\log_b a > k$, then

$$\mathcal{O}(n^{\log_b a}).$$

Case 2: if $\log_b a = k$, then

- if $p > -1$,
$$\mathcal{O}(n^k \log^{p+1} n),$$
- if $p = -1$,
$$\mathcal{O}(n^k \log \log n),$$
- if $p < -1$,
$$\mathcal{O}(n^k).$$

Case 3: if $\log_b a < k$, then

- if $p \geq 0$,
$$\mathcal{O}(n^k \log^p n),$$
- if $p < 0$,
$$\mathcal{O}(n^k).$$