

CSCI3230 Fundamentals of Artificial Intelligence

Ryan Chan

December 1, 2025

Abstract

This is a note for **CSCI3230 Fundamentals of Artificial Intelligence**.

Contents are adapted from the lecture notes of CSCI3230, prepared by [Qi Dou](#), as well as some online resources.

This note is intended solely as a study aid. While I have done my best to ensure the accuracy of the content, I do not take responsibility for any errors or inaccuracies that may be present. Please use the material thoughtfully and at your own discretion.

If you believe any part of this content infringes on copyright, feel free to contact me, and I will address it promptly.

Mistakes might be found. So please feel free to point out any mistakes.

Contents

1	Basis of AI	2
1.1	Introduction	2
1.2	Model Performance	3
1.3	Types of Learning	4
2	Regression	5
2.1	Linear Regression	5
2.2	Logistic Regression	11
3	Support Vector Machine	12
4	Clustering Algorithms	13
5	Neural Network Basics	14
6	Convolutional Neural Networks	15
7	Transformers	16
8	Uninformed Search	17
9	Informed Search	18
10	AI Applications	19

Chapter 1

Basis of AI

1.1 Introduction

Artificial intelligence (AI) is intelligence demonstrated by machines, as opposed to the natural intelligence displayed by humans or animals.¹

It can be distinguished into **Artificial Narrow Intelligence** and **Artificial General Intelligence**. The latter refers to the ability of an intelligent agent to understand or learn any intellectual task that a human being can. Machine learning, as often mentioned, is a branch of AI, while deep learning is a specific subset of machine learning.

Before delving into details, we consider how humans learn, since AI is basically simulating the way people learn.

For humans, we acquire or learn some skills by accumulating experience from observations. This is basically the same for computers, except that the observations are now in the form of data, which can vary from text documents to sound recordings.

The skill of a machine can be illustrated by prediction, which is basically what AI does—makes predictions. Then, how can data be learned? There are learning algorithms that train the machine on how to make predictions; we construct a model corresponding to such an algorithm and then use the data to train the model.

Think of it this way: a learning algorithm is the functionality of the brain, where the model is the container, i.e., the brain itself. We have observations (training data), and then, using these observations, we make predictions based on new information (test data) — for example, the answer to a math question, the answer to someone else's question, etc.

Then, you observe that there are two sets of data: one is the training set and the other is the test set. Their usage is just like their names suggest. For example, for labeled data, we might randomly split them into 80% training data and 20% test data. Notice that a larger test set will give a more accurate assessment of performance.

We sometimes also use a validation data set, which is a set of examples used to tune the *hyperparameters* of the model.

In short, we have an original labeled data set, and we either split it into two sets, i.e., training and test sets, or we split it into three sets with a validation set used to tune the model.

If there is a small-scale dataset, we can use **cross-validation**. We shuffle the dataset randomly and split it into k groups. For each unique group, take that group as a hold-out or test dataset, and use the remaining groups as the training dataset. Then, summarize the performance of the model using all the evaluation scores from each fold.

With all these methods, the question remains: how do we evaluate the performance of a model given a dataset?

¹Here are some funny facts I would like to share, which are not necessary. You can directly go to [1.2](#).

1.2 Model Performance

To evaluate the model performance, we can compare it with previous models if they exist. There are also some metrics that can be used, which are based on the confusion matrix.

Assume we have a binary classifier that can classify the data into positive and negative samples. Then we have

		Actual	
		Positive	Negative
Predicted	Positive	True Positive (TP)	False Positive (FP)
	Negative	False Negative (FN)	True Negative (TN)

Table 1.1: Standard Binary Confusion Matrix Layout

Here TP and TN show correct prediction, whereas FP and FN show incorrect prediction.

Below might provide a more intuitive illustration.²



Figure 1.1: Pregnancy Confusion Matrix

We define **Accuracy** as the fraction of correct prediction,

$$\frac{TP + TN}{TP + FP + TN + FN}.$$

Precision is the accuracy of the positive predictions:

$$\frac{TP}{TP + FP}.$$

Notice that high precision means low ‘false alarm rate’.

Lastly, **recall** (or **sensitivity**) is the accuracy for the positive class,

$$\frac{TP}{TP + FN}.$$

High recall means low possibility of missing many positives.

We use a confusion matrix to visualize the performance of an algorithm. Each row of the matrix represents the instances in an actual class, while each column represents the instances in a predicted class, or vice versa.

²<https://codefinity.com/courses/v2/b71ff7ac-3932-41d2-a4d8-060e24b00129/6b1d35ef-cd16-405c-8202-ef0f3d8e10c5/be441b8d-e533-4c4d-8e34-817ea61b6fa0>

1.3 Types of Learning

There are three types of learning.

Supervised Learning: data is labeled such that the machine can learn from the characters. This is used to predict outcomes or the future. It can be seen as a function that maps an input to an output based on example input–output pairs.

Unsupervised Learning: data is in the form of unlabeled samples, where the machine needs to learn on its own and find hidden structure.

Reinforcement Learning: this can be considered as a reward system, where it deals with the way intelligent agents take actions in order to maximize their cumulative reward.

Chapter 2

Regression

Regression is a statistical process for estimating the relation between dependent and independent variables. Consider it as a program, where the input is independent variables and the output is dependent variables.

2.1 Linear Regression

Linear regression is a kind of regression model, which hypothesizes that the relation between dependent variables and independent variables is linear. Mathematically, a function $f(x)$ is linear iff. $f(u + v) = f(u) + f(v)$ and $f(cu) = cf(u)$.

2.1.1 Formulation

Consider intelligent machines as functions with input and output. Then we have

$$f : X \rightarrow Y$$

where $X \subset \mathbb{R}^n$ is the domain of input and $Y \subset \mathbb{R}$ is the domain of output. The goal of the learning algorithm is to take the given data as training examples, then try to find a general mapping $\hat{f} : X \rightarrow Y$ such that it is close to the true f , i.e. $\hat{f} \approx f$.

As we plot the given data on a graph, if the data points are distributed along a straight line, we can assume the relation between input X and output Y is linear for simplification.

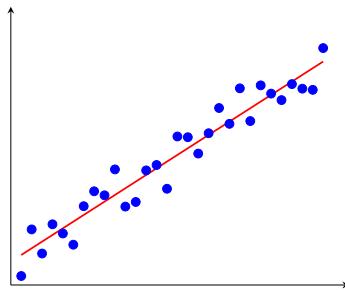


Figure 2.1: Linear Regression

For a univariate linear function in a 1-dimensional feature space, we have

$$\hat{y} = \hat{f}_{\theta_0, \theta_1}(x) = \theta_0 + \theta_1 x.$$

If we can find the optimal θ_0 and θ_1 , i.e., the relation parameters, then we can predict the output with higher precision.

When there are more than one factors, we consider an n -dimensional feature space. The multivariate linear function is

$$\hat{y} = \hat{f}_{\theta_0, \theta_1, \dots, \theta_n}(x_1, \dots, x_n) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n.$$

We can also write it in vector form as

$$\hat{y} = \hat{f}_{\theta_0, \Theta}(X) = X^T \Theta + \theta_0,$$

where

$$\Theta = \begin{pmatrix} \theta_1 \\ \vdots \\ \theta_n \end{pmatrix}, \quad X = \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}.$$

Remark. θ_0 and Θ are parameters that should be learned from the training data.

There might be more than one sample being studied. Suppose we have m samples, and

$$X^{(1)} = \begin{pmatrix} x_1^{(1)} \\ \vdots \\ x_n^{(1)} \end{pmatrix}, \quad X^{(2)} = \begin{pmatrix} x_1^{(2)} \\ \vdots \\ x_n^{(2)} \end{pmatrix}, \dots, X^{(m)} = \begin{pmatrix} x_1^{(m)} \\ \vdots \\ x_n^{(m)} \end{pmatrix},$$

with respective labels $y^{(1)}, y^{(2)}, \dots, y^{(m)}$.

We denote $\mathbf{X} \in \mathbb{R}^{m \times n}$ as the data matrix, where rows represent samples and columns represent features:

$$\mathbf{X} = \begin{pmatrix} X^{(1)^T} \\ \vdots \\ X^{(m)^T} \end{pmatrix} = \begin{pmatrix} x_1^{(1)} & \cdots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_1^{(m)} & \cdots & x_n^{(m)} \end{pmatrix}.$$

Then we have a more general function for linear regression with n features and m samples:

$$\hat{Y} = \hat{f}_{\theta_0, \Theta}(\mathbf{X}) = \mathbf{X}\Theta + \theta_0,$$

where

$$\Theta = \begin{pmatrix} \theta_1 \\ \vdots \\ \theta_n \end{pmatrix}, \quad \hat{Y} = \begin{pmatrix} \hat{y}^{(1)} \\ \vdots \\ \hat{y}^{(m)} \end{pmatrix}.$$

Remark. For $\mathbf{X}\Theta$, consider it as combining the features for each sample:

$$\mathbf{X}\Theta = \theta_1 \begin{pmatrix} x_1^{(1)} \\ \vdots \\ x_1^{(m)} \end{pmatrix} + \theta_2 \begin{pmatrix} x_2^{(1)} \\ \vdots \\ x_2^{(m)} \end{pmatrix} + \dots + \theta_n \begin{pmatrix} x_n^{(1)} \\ \vdots \\ x_n^{(m)} \end{pmatrix}.$$

For simplification, we define

$$\mathbf{X} = \begin{pmatrix} 1 & x_1^{(1)} & \cdots & x_n^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(m)} & \cdots & x_n^{(m)} \end{pmatrix}, \quad \Theta = \begin{pmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{pmatrix},$$

then

$$\hat{Y} = \hat{f}_{\Theta}(\mathbf{X}) = \mathbf{X}\Theta.$$

Note (Geometry of Linear Regression). Data points $\{(x_1^{(1)}, \dots, x_n^{(1)}, y^{(1)})\}, \dots, \{(x_1^{(m)}, \dots, x_n^{(m)}, y^{(m)})\}$ occupy a $(n+1)$ -dimensional space. The fitted line is actually an n -dimensional hyperplane in this space.

Since all data points are supposed to be sampled from the same distribution, once we find the optimal Θ , we obtain $\hat{f}_\Theta \approx f$. The next question is: how do we find the optimal Θ ? We first focus on fitting the training data.

2.1.2 Optimization

Intuitively, we need to measure and minimize the distance between the estimated values $\hat{f}_\Theta(X^{(i)})$ and the true training labels $y^{(i)}$.

We denote $J(\Theta)$ as the objective function (or cost function) to measure this distance. The goal is to find Θ such that $J(\Theta)$ is minimized.

We define $J(\Theta)$ as the **Residual Sum of Squares (RSS)**:

$$J(\Theta) = \sum_{i=1}^m (\hat{f}_\Theta(X^{(i)}) - y^{(i)})^2 = \|\hat{f}_\Theta(\mathbf{X}) - Y\|_2^2.$$

The method that approximates the solution by minimizing the RSS is called **ordinary least squares**.

Remark. The *residual* is the difference between the estimated value and the corresponding training label.

Consider the univariate linear regression problem. Its residual sum of squares (RSS) is

$$J(\theta_0, \theta_1) = \sum_{i=1}^m (\hat{f}_{\theta_0, \theta_1}(x^{(i)}) - y^{(i)})^2 = \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2.$$

Then we can rewrite $J(\theta_0, \theta_1)$ as

$$\begin{aligned} J(\theta_0, \theta_1) &= \sum_{i=1}^m (x^{(i)} - \bar{x})^2 \left(\theta_1 - \frac{\sum_{i=1}^m (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^m (x^{(i)} - \bar{x})^2} \right)^2 + m(\theta_0 - (\bar{y} - \theta_1 \bar{x}))^2 \\ &\quad + \sum_{i=1}^m (y^{(i)} - \bar{y})^2 - \frac{\left(\sum_{i=1}^m (x^{(i)} - \bar{x})(y^{(i)} - \bar{y}) \right)^2}{\sum_{i=1}^m (x^{(i)} - \bar{x})^2}. \end{aligned}$$

Remark. Only the first two terms depend on θ_0 and θ_1 . Therefore, to minimize $J(\theta_0, \theta_1)$, it suffices to set these two terms to zero.

Thus, we solve

$$\begin{cases} 0 = \sum_{i=1}^m (x^{(i)} - \bar{x})^2 \left(\theta_1 - \frac{\sum_{i=1}^m (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^m (x^{(i)} - \bar{x})^2} \right)^2, \\ 0 = m(\theta_0 - (\bar{y} - \theta_1 \bar{x}))^2, \end{cases}$$

which gives the closed-form solution

$$\begin{cases} \theta_1 = \frac{\sum_{i=1}^m (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_{i=1}^m (x^{(i)} - \bar{x})^2}, \\ \theta_0 = \bar{y} - \theta_1 \bar{x}. \end{cases}$$

Usually, we use optimization methods to find Θ^* that minimizes the objective function. Optimization is the process of selecting the best element with regard to some criterion. We can write it as

$$\arg \min_x g(x) \quad \text{s.t. some conditions here,}$$

where $g(x)$ is the target to minimize and x is the element to search.

For ordinary linear regression, we can formulate the optimization problem as

$$\hat{\Theta} = \Theta^* = \arg \min_{\Theta} J(\Theta).$$

Since the RSS function is convex, i.e., the global minimum is also a local minimum, we can find the global minimum by solving

$$J'(\Theta^*) = 0 \quad \text{and} \quad J''(\Theta^*) > 0.$$

We have

$$\begin{aligned} J(\Theta) &= \|\hat{f}_{\Theta}(\mathbf{X}) - Y\|_2^2 = (\mathbf{X}\Theta - Y)^T(\mathbf{X}\Theta - Y) \\ &= \Theta^T \mathbf{X}^T \mathbf{X}\Theta - Y^T \mathbf{X}\Theta - \Theta^T \mathbf{X}^T Y - Y^T Y. \end{aligned}$$

Then, the derivatives are

$$\begin{aligned} \frac{\partial J(\Theta)}{\partial \Theta} &= 2\mathbf{X}^T \mathbf{X}\Theta - \mathbf{X}^T Y - \mathbf{X}^T Y \\ &= 2\mathbf{X}^T(\mathbf{X}\Theta - Y) = 0, \\ \frac{\partial^2 J(\Theta)}{\partial \Theta^2} &= 2\mathbf{X}^T \mathbf{X} > 0. \end{aligned}$$

Finally, we obtain the closed-form solution

$$\mathbf{X}^T \mathbf{X}\Theta = \mathbf{X}^T Y \implies \hat{\Theta} = \Theta^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T Y.$$

Intuition. All in all, the goal is to find the best Θ , i.e., the parameters, since we want the model to capture most of the data and provide more precise predictions. For the case with only one feature, the derived formula is sufficient. However, for multiple features and samples, the calculation becomes cumbersome.

Thus, we consider the RSS function in another way. Since it is convex by nature, we can use its derivative to find the optimal Θ efficiently.

2.1.3 Shrinkage Methods

There are noises from observation. Therefore, we have $y = f(X) + \epsilon$, where $\mathbb{E}(\epsilon) = 0$ and $\text{Var}(\epsilon) = \sigma^2$. Noise could lead to error, so for any fixed X with label y the expected prediction error at X is

$$\text{EPE}(X) = \mathbb{E}((y - \hat{f}(X))^2) = \text{Bias}(\hat{f}(X))^2 + \text{Var}(\hat{f}(X)) + \sigma^2$$

where

$$\text{Bias}(\hat{f}(X)) = \mathbb{E}(\hat{f}(X)) - f(X) \quad \text{Var}(\hat{f}(X)) = \mathbb{E}((\hat{f}(X) - \mathbb{E}[\hat{f}(X)])^2)$$

The expected prediction error is composed of bias, variance, and irreducible error. We can achieve a considerable EPE if both bias and variance are minimized simultaneously (see Figure 2.2¹).

As we repeat the training on randomly sampled data, if the model perfectly fits the training data, the prediction bias is low while variance is very high, since the model varies among different training datasets. Yet, if the model is constant among different training datasets, the prediction variance will be zero but with high bias. This means that in general, low variance causes high bias, and vice versa.

Thus, we need to make a trade-off between minimizing bias and variance.

¹<http://scott.fortmann-roe.com/docs/BiasVariance.html>

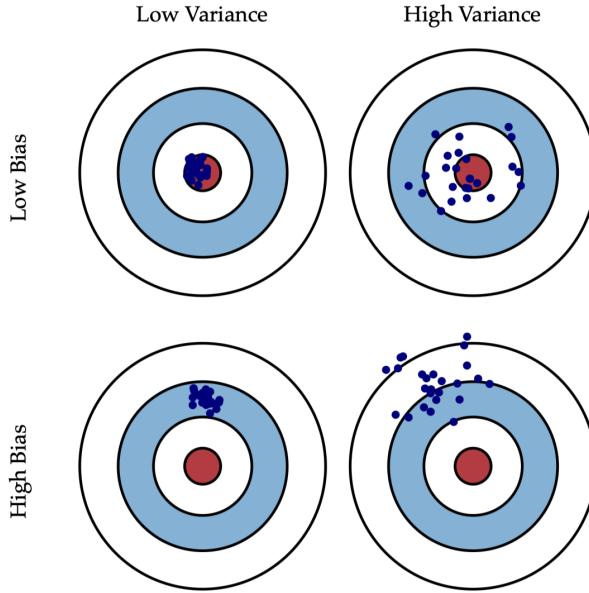


Figure 2.2: Graphical illustration of bias and variance

Over-fitting refers to the case where we have low bias but high variance, which usually occurs when the model has high complexity, i.e., many features used in training and prediction.

Under-fitting refers to the case where we have high bias but low variance, which occurs when the model complexity is too low.

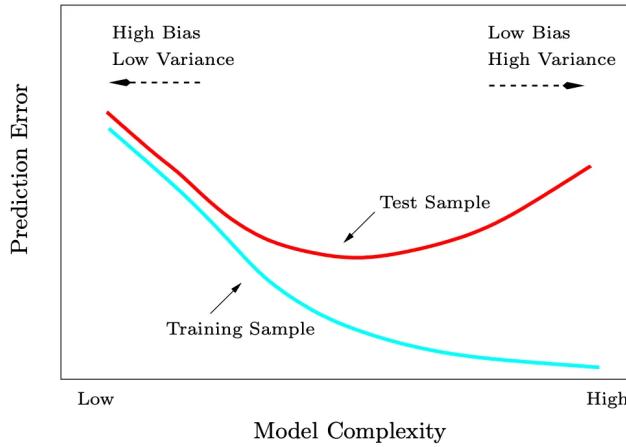


Figure 2.3: Model Complexity

Ordinary least squares is prone to producing a model with low bias but high variance. To improve the EPE, we can tune the model complexity. Shrinkage methods are a set of techniques for automatically controlling and reducing model complexity.

Ridge Regression

When there are many correlated features in a linear regression model, the estimated coefficients can become poorly determined and exhibit high variance. For example, a very large positive coefficient on one variable can be offset by a similarly large negative coefficient on a correlated variable. The ridge penalty regularizes the coefficients by adding a cost for large values, so that the optimization function discourages excessively large coefficients. As a result, the problem of unstable estimates due to multicollinearity is alleviated.

As we have

$$\hat{Y} = \hat{f}_{\theta_0, \Theta}(\mathbf{X}) = \mathbf{X}\Theta + \theta_0,$$

the ordinary least squares problem can be formulated as the following optimization:

$$\hat{\Theta} = \arg \min_{\Theta} \|\mathbf{X}\Theta + \theta_0 - Y\|_2^2.$$

Ridge regression aims to shrink the parameters by imposing a penalty on the squared magnitude of the coefficients:

$$\hat{\Theta}^{ridge} = \arg \min_{\Theta} \|\mathbf{X}\Theta + \theta_0 - Y\|_2^2 + \lambda \|\Theta\|_2^2.$$

The penalty term (also known as the L2 norm) is defined as

$$\|\Theta\|_2^2 = \sum_{i=1}^n \theta_i^2.$$

Note that θ_0 is not included in the penalty term, and λ is a user-specified hyperparameter.

In the figures below, the red '+' markers represent points on the true relationship $f(X)$ (shown as blue lines). Ordinary least squares fits the training data very well but exhibits high variability across different training sets. In contrast, ridge regression reduces the variability of the estimated coefficients, leading to more stable models, but may not fit the training data perfectly.

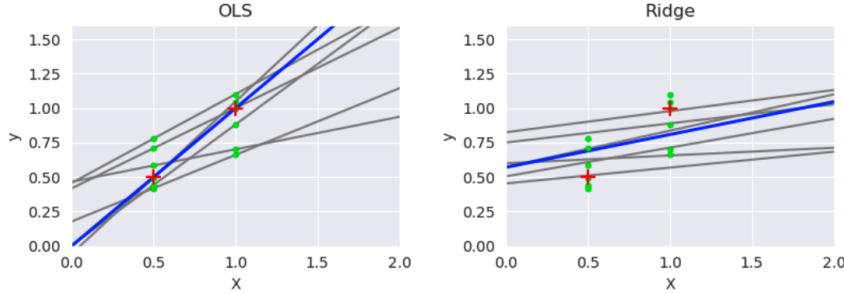


Figure 2.4: OLS v.s. Ridge

Intuition. In a test, each feature represents a factor affecting your score: hours studied, sleep, coffee, etc. In OLS, you try to perfectly fit the past test scores. Some factors may receive very large weights to match the training data exactly.

However, if a future test changes slightly (due to noise), your predictions fluctuate a lot, resulting in high variance. Ridge regression addresses this by discouraging any factor from having too much influence. If a parameter becomes too large, it incurs a penalty.

As a result, the optimization shrinks Θ toward zero.

Lasso Regression

Lasso regression is also a shrinkage method, similar to ridge regression, but differs in the penalty term:

$$\hat{\Theta}^{lasso} = \arg \min_{\Theta} \|\mathbf{X}\Theta + \theta_0 - Y\|_2^2 + \lambda \|\Theta\|_1.$$

The penalty term (also known as the L1 norm) is defined as

$$\|\Theta\|_1 = \sum_{i=1}^n |\theta_i|,$$

and is designed to encourage sparsity in the parameters. Sparsity means that many elements in Θ are exactly zero, effectively reducing the number of features used and therefore decreasing model complexity.

Similarly to ridge regression, Lasso yields models with less variance, but it is less likely to perfectly fit every data point.

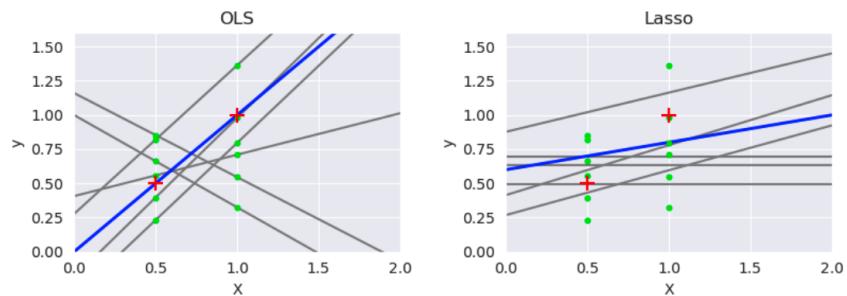


Figure 2.5: OLS v.s. Ridge

2.2 Logistic Regression

Chapter 3

Support Vector Machine

Chapter 4

Clustering Algorithms

Chapter 5

Neural Network Basics

Chapter 6

Convolutional Neural Networks

Chapter 7

Transformers

Chapter 8

Uniformed Search

Chapter 9

Informed Search

Chapter 10

AI Applications